

---

# Project Synopsis

## Autonomous ML Model Monitoring & Drift Management Platform

---

### Cover Information:

- **Project Title:** Autonomous ML Model Monitoring & Drift Management Platform
  - **Project Type:** Mini Project
  - **Team Name/ID:** Team 76
  - **Course:** B. Tech CSE – AIML
  - **Institute:** GLA University, Mathura
  - **Date:** 08 Feb 2026
- 

## 1. Overview

### 1.1 Problem Statement

Machine Learning models deployed in production environments are subject to performance degradation over time due to changes in data distribution, business environment, and user behavior. This degradation, known as **Model Drift**, can cause significant operational and financial losses.

Current monitoring approaches are either:

- Manual
- Periodic
- Reactive
- Dependent on human supervision

There is no fully autonomous system that detects drift and retrains models automatically in an academic-friendly scalable format.

---

### 1.2 Goal

To develop an enterprise-grade autonomous platform that:

- Continuously monitors ML models
  - Detects data, concept, and prediction drift
  - Automatically triggers retraining
- 
- Deploys updated model versions
  - Provides centralized dashboard monitoring
- 

### 1.3 Non-Goals

- Reinforcement learning adaptive systems
  - Federated distributed training
  - Cross-region fault tolerance
  - Automated compliance certification
- 

### 1.4 Value Proposition

| Feature         | Traditional Monitoring | Proposed Platform              |
|-----------------|------------------------|--------------------------------|
| Drift Detection | Manual                 | Automated                      |
| Retraining      | Human-triggered        | Autonomous                     |
| Monitoring      | Periodic               | Continuous                     |
| Deployment      | Manual                 | Version-controlled auto-deploy |
| Cost            | High enterprise        | Academic scalable              |

---

## 2. Scope and Control

### 2.1 In-Scope

- Model registration
  - Data monitoring
  - Statistical drift detection
  - Retraining pipeline
  - Alert system
  - Dashboard
- 

### 2.2 Out-of-Scope

- Multi-asset real-time GPU orchestration
  - Blockchain audit storage
  - Mobile application
- 

### 2.3 Assumptions

| Assumption              | Description                |
|-------------------------|----------------------------|
| Training Data Available | Historical baseline exists |
| Labels Accessible       | For retraining validation  |
| Compute Available       | Cloud/Local server         |

| Assumption               | Description      |
|--------------------------|------------------|
| API Integration Possible | For model access |

---



---

## 2.4 Constraints

| Constraint            | Impact                |
|-----------------------|-----------------------|
| Academic timeline     | Limited iteration     |
| Limited cloud credits | Resource optimization |
| Compute limitations   | Lightweight models    |
| API rate limits       | Batch processing      |

---

## 2.5 Acceptance Criteria

| Scenario           | Condition            | Expected Outcome     |
|--------------------|----------------------|----------------------|
| Drift Occurs       | PSI > 0.25           | Retraining triggered |
| Accuracy Drop      | Accuracy < threshold | Alert generated      |
| Retraining Success | Accuracy improves    | New version deployed |
| System Health      | API response         | < 1000 ms            |

---

---

### 3. Stakeholders and RACI

| 3Activity                | Responsible (R) | Accountable (A) | Consulted (C) | Informed (I) |
|--------------------------|-----------------|-----------------|---------------|--------------|
| Requirement Analysis     | Vaishnavi       | Anjali          | Mentor        | Team         |
| Architecture Design      | Abhay           | Anjali          | Mentor        | Team         |
| Drift Algorithm Research | Anjali          | Anjali          | Mentor        | Team         |
| Backend Development      | Abhay           | Abhay           | Team          | Mentor       |
| Dashboard Development    | Vaishnavi       | Abhay           | Team          | Mentor       |
| Testing & QA             | Vaishnavi       | Anjali          | Mentor        | Team         |
| Final Release            | Anjali          | Anjali          | Mentor        | Department   |

---

### 4. Team and Roles

| Member           | Role             | Responsibilities                                       | Key Skills      | Availability |
|------------------|------------------|--------------------------------------------------------|-----------------|--------------|
| Anjali Kumari    | ML Lead          | Drift detection research, validation, retraining logic | ML, Statistics  | 10 hrs/week  |
| Abhay Tarkar     | Backend & DevOps | APIs, DB design, Docker deployment                     | Python, FastAPI | 12 hrs/week  |
| Vaishnavi Pandey | Frontend & QA    | Dashboard UI, Testing, Documentation                   | React, Testing  | 8 hrs/week   |

---

## 5. Week-wise Plan and Assignments

| Week | Milestone               | Anjali              | Abhay           | Vaishnavi        | Deliverables   |
|------|-------------------------|---------------------|-----------------|------------------|----------------|
| 1    | Requirements Freeze     | Scope finalization  | API planning    | Documentation    | SRS Draft      |
| 2    | Baseline Model          | Dataset prep        | DB schema       | UI wireframe     | Baseline model |
| 3    | Drift Simulation        | Statistical testing | API scaffolding | Test cases       | Drift report   |
| 4    | Backend Integration     | Validation          | API dev         | UI skeleton      | Working API    |
| 5    | Retraining Engine       | Threshold tuning    | Pipeline dev    | Dashboard charts | Auto retrain   |
| 6    | Dashboard Finalization  | Metrics analysis    | Deployment      | UI polish        | Demo build     |
| 7    | Testing & QA            | Evaluation          | Bug fixes       | E2E testing      | Test report    |
| 8    | Release & Documentation | Final validation    | Deployment      | Final docs       | Submission     |

## 6. Users and UX


### 6.1 Personas

| Persona          | Need              | Pain Point         |
|------------------|-------------------|--------------------|
| ML Engineer      | Auto monitoring   | Silent failures    |
| Data Scientist   | Drift analytics   | Manual checks      |
| Enterprise Admin | Central dashboard | Fragmented systems |

### 6.2 Top User Journey

User → Login → Register Model → Monitor Dashboard → Receive Drift Alert → Auto Retrain → Version Update

### 6.3 Use Case Diagram



The Use Case Diagram represents the functional interaction between system users and the Autonomous ML Monitoring System. It identifies the primary actors and the operations they can perform within the system boundary.

There are two main actors in the system:

### 1. ML Engineer

The ML Engineer is responsible for managing machine learning models and monitoring their performance in production. This actor can:

- Register a new model
- Upload training data
- Monitor model performance
- View drift metrics
- View model versions
- Receive alerts regarding performance degradation

The ML Engineer primarily interacts with monitoring and analytical functionalities of the system.

---

### 2. Admin

The Admin oversees system governance and operational control. The Admin can:

- Manage users
- Monitor model performance
- View drift metrics
- Trigger retraining manually if required

The Admin ensures system-level control and operational integrity.

---

### Core Functional Flow

The diagram highlights important <> relationships that represent internal system dependencies:

- **Monitor Model Performance** includes **Detect Drift**  
This means drift detection is automatically performed whenever model performance is monitored.
-



- **Detect Drift** includes **Trigger Retraining**

If drift exceeds predefined thresholds, the system automatically initiates retraining.

These include relationships show that retraining is not a standalone action but is triggered as a consequence of drift detection.

---

### System Boundary

All use cases are enclosed within the system boundary labeled

**"Autonomous ML Monitoring System"**, indicating that these operations are handled internally by the platform.

Actors remain outside the boundary to represent external interaction with the system.

---

### Functional Significance

The Use Case Diagram demonstrates that the system:

- Supports both operational users (ML Engineers) and administrative users
- Provides automated internal processes (drift detection and retraining)
- Ensures controlled interaction between users and core ML lifecycle management

This structured representation validates that the system satisfies functional requirements and supports autonomous ML model management.

---

## 7. Market and Competitors

| Competitor   | Target            | Key Features    | Weakness      | Our Differentiator    |
|--------------|-------------------|-----------------|---------------|-----------------------|
| Evidently AI | ML Engineers      | Drift detection | No automation | Autonomous retraining |
| Arize AI     | Enterprises       | Observability   | Expensive     | Academic scalable     |
| Fiddler AI   | AI explainability | Monitoring      | Complex setup | Simplified system     |

---

## 8. Objectives and Success Metrics

| Objective                | KPI                   | Target               |
|--------------------------|-----------------------|----------------------|
| Drift Detection Accuracy | Precision             | $\geq 90\%$          |
| Accuracy Recovery        | Post-retrain accuracy | $\geq 95\%$ baseline |
| Latency                  | p95                   | $\leq 1000$ ms       |
| Availability             | Uptime                | $\geq 99\%$          |
| Error Rate               | 5xx %                 | $\leq 1\%$           |

---

## 9. Key Features

| Feature            | Description          | Priority | Dependencies | Acceptance Criteria  |
|--------------------|----------------------|----------|--------------|----------------------|
| Model Registration | Upload & store model | Must     | Auth         | Model saved          |
| Drift Detection    | PSI, KS, KL tests    | Must     | Data         | Drift flagged        |
| Auto Retraining    | Pipeline trigger     | Must     | Drift engine | New version deployed |
| Alert System       | Email notifications  | Should   | SMTP         | Alert within 1 min   |
| Dashboard          | Visual metrics       | Must     | API          | Live monitoring      |

## 10. Architecture

### 10.1 High-Level Architecture

Client (React SPA)

↓

Backend API (FastAPI)

↓

Monitoring Engine

↓


Drift Detection Module

↓

Retraining Orchestrator

↓

Model Registry + Database



**Figure 10.1 :** High-Level Architecture of Autonomous ML Monitoring Platform

The High-Level Architecture diagram illustrates the layered structure and overall system flow of the Autonomous ML Model Monitoring & Drift Management Platform. The system is designed using a modular, layered architecture to ensure scalability, maintainability, and clear separation of concerns.

---

## **1 Client Layer**

The Client Layer consists of the Web Dashboard (React) and the Admin Interface. This layer provides a user-friendly interface through which ML Engineers and Administrators interact with the system.

Users can:

- Register models
- Monitor performance metrics
- View drift reports
- Trigger retraining (if authorized)
- Receive alerts

All user requests are routed securely to the Backend API layer.

---

## **2 API Layer**

The API Layer acts as the communication bridge between the frontend and backend services. It includes:

- Backend API (FastAPI)
- Authentication Service (JWT-based)

The Backend API handles:

- Request validation
- Model registration
- Metric retrieval
- Drift check invocation
- Retraining triggers

The Authentication Service ensures secure access control using JWT tokens and role-based permissions.

---

---

### **Core Monitoring Layer**

This is the central intelligence layer of the system.

It contains:

- Monitoring Engine
- Drift Detection Module
- Performance Evaluator

The Monitoring Engine continuously receives production data streams and interacts with the deployed ML model to generate predictions.

The Drift Detection Module performs statistical analysis using:

- Population Stability Index (PSI)
- Kolmogorov-Smirnov Test (KS Test)
- KL Divergence

If the computed drift score exceeds predefined thresholds, the system identifies a drift event.

The Performance Evaluator tracks accuracy, precision, recall, and other metrics over time.

---

### **Automation Layer**

Once drift is detected, the Automation Layer takes control.

It includes:

- Retraining Orchestrator
- Model Validation Module
- Version Controller

The Retraining Orchestrator automatically initiates model retraining using updated datasets. The Model Validation Module evaluates retrained model performance.

If validation metrics meet the acceptance criteria, the Version Controller updates the active model version and archives the previous one.

This layer enables the system to be self-healing and autonomous.

---

## **5 Data & Storage Layer**

This layer ensures persistence and traceability.

It includes:

- Model Registry
- PostgreSQL Database
- Log Storage

The Model Registry maintains version history and metadata.

The database stores users, models, drift reports, and alerts.

Log Storage ensures auditability and trace tracking for compliance and debugging.

---

## **6 External Components**

External systems interacting with the platform include:

- Deployed ML Model
- Production Data Stream
- Email/Alert Service

Production data flows into the Monitoring Engine.

Drift alerts are sent via the Email/Alert Service.

The deployed ML model continuously interacts with the monitoring engine for prediction generation.

---

## **Architectural Significance**

This layered architecture ensures:

- Clear separation of concerns
- Scalability
- Fault isolation
- Secure access control
- Automated lifecycle management

By separating monitoring, automation, and storage layers, the system achieves enterprise-grade modularity while maintaining academic project feasibility.

---

The architecture demonstrates a complete ML lifecycle management system capable of autonomous drift detection and correction.

---

## 10.2 API Spec Snapshot

| Endpoint            | Method | Auth | Purpose         | Response     |
|---------------------|--------|------|-----------------|--------------|
| /api/auth/login     | POST   | —    | Login user      | JWT          |
| /api/model/register | POST   | JWT  | Register model  | ModelID      |
| /api/drift/check    | GET    | JWT  | Check drift     | Drift score  |
| /api/retrain        | POST   | JWT  | Trigger retrain | VersionID    |
| /api/metrics        | GET    | JWT  | Fetch stats     | JSON metrics |

---

## 11. Data Design

### 11.1 Core Entities

- User
  - Model
  - ModelVersion
  - DriftReport
  - AlertLog
-

## 11.2 Data Dictionary

| Entity       | Field       | Type   | Null? | Description      |
|--------------|-------------|--------|-------|------------------|
| User         | id          | UUID   | No    | Primary Key      |
| User         | email       | String | No    | Unique           |
| Model        | id          | UUID   | No    | Model identifier |
| DriftReport  | drift_score | Float  | No    | PSI score        |
| ModelVersion | accuracy    | Float  | No    | Validation score |


---

## 11.3 Privacy & Backup

- JWT authentication
  - Encrypted storage
  - Weekly backups
  - RTO: 4 hours
  - RPO: 24 hours
-



## 11.4 Entity-Relationship Diagram



**Figure 11.4: Entity Relationship Diagram (ERD) for Autonomous ML Model Monitoring & Drift Management Platform**

### Overview

This diagram illustrates the relational database schema designed to support the lifecycle management and monitoring of machine learning models. It defines five core entities and their logical associations, utilizing standard Chen notation with rectangular entities and oval attributes.

### Detailed Description

#### 1. User Management:

The **User** entity represents system actors (e.g., Data Scientists, DevOps). It holds a **One-to-Many (1:N)** relationship with the **Model** entity, signifying that a single registered user can own and manage multiple machine learning models within the platform.

#### 2. Model Versioning:

The **Model** entity serves as the central node. It maintains a **One-to-Many (1:N)** relationship with **ModelVersion**. This structure allows the system to track the historical evolution of a model (e.g., distinct version numbers and accuracy metrics) while maintaining a reference to the parent model ID.

3. **Drift Detection:**

The system monitors model performance via the **DriftReport** entity. A **One-to-Many (1:N)** relationship exists between **Model** and **DriftReport**, allowing a single model to generate a history of drift assessments over time (tracking attributes like `drift_score` and `drift_type`).

4. **Alerting System:**

When drift is detected, the system triggers notifications. The **DriftReport** entity has a **One-to-Many (1:N)** relationship with **AlertLog**. This ensures that a single drift event can generate multiple alert records (e.g., distinct logs for email, dashboard, or SMS notifications).

**Key Notation**

- **PK (Primary Key):** Unique identifiers for each entity (e.g., `model_id`, `report_id`).
- **FK (Foreign Key):** Attributes used to link entities and enforce referential integrity (e.g., `user_id` in the **Model** entity).

12. **Drift Detection Algorithms**

| Algorithm     | Purpose                  | Threshold                  |
|---------------|--------------------------|----------------------------|
| PSI           | Distribution shift       | >0.25 severe               |
| KS Test       | Statistical difference   | $p < 0.05$                 |
| KL Divergence | Probabilistic divergence | High value indicates drift |
| Wasserstein   | Distance metric          | Threshold-based            |

13. **Non-Functional Requirements**

| Metric       | SLI    | Target (SLO) | Measurement     |
|--------------|--------|--------------|-----------------|
| Availability | Uptime | $\geq 99\%$  | Monitoring logs |

| Metric     | SLI           | Target (SLO) | Measurement |
|------------|---------------|--------------|-------------|
| Latency    | p95           | ≤ 1000 ms    | APM         |
| Error Rate | 5xx %         | ≤ 1%         | Logs        |
| Security   | Critical CVEs | 0            | Scanner     |

---

## 14. Test Plan

| Area    | Type        | Tool       | Owner     | Coverage      | Exit Criteria |
|---------|-------------|------------|-----------|---------------|---------------|
| Backend | Unit        | Pytest     | Abhay     | 70%           | No P1 defects |
| UI      | E2E         | Playwright | Vaishnavi | 60%           | ≥ 95% pass    |
| API     | Integration | Postman    | Anjali    | All scenarios | Critical pass |

---

## 15. Security & Compliance

| Asset       | Threat              | Impact | Mitigation            |
|-------------|---------------------|--------|-----------------------|
| Auth Tokens | Theft               | High   | HTTPS, TTL            |
| User Data   | SQL Injection       | High   | Parameterized queries |
| Models      | Unauthorized access | Medium | RBAC                  |

---

## 16. Risks and Mitigation

| Risk               | Probability | Impact | Score | Mitigation        |
|--------------------|-------------|--------|-------|-------------------|
| False Drift Alerts | Medium      | Medium | 9     | Threshold tuning  |
| API Downtime       | Low         | High   | 8     | Backup monitoring |
| Schedule Delay     | Medium      | High   | 12    | Weekly review     |

---

## 17. Research and Evaluation

Existing platforms provide monitoring but lack autonomous retraining.

Experimental Results:

| Stage            | Accuracy |
|------------------|----------|
| Baseline         | 91%      |
| After Drift      | 78%      |
| After Retraining | 90%      |

System achieved  $\geq 92\%$  drift detection precision.

---

---

## 18. Appendices

### Glossary

- Drift – Change in data distribution
- PSI – Population Stability Index
- SLO – Service Level Objective
- p95 – 95th percentile latency

### References

- Scikit-learn documentation
  - Google MLOps Whitepaper
  - IEEE papers on Concept Drift
  - Evidently AI documentation
-